

Week 3: Data Visualization

`{ggplot2}`


Joe Nese

University of Oregon

Fall 2022

Housekeeping

- Homework 1 due today
- Groups due today
 - I need emails from everyone by 5pm (if you have not already emailed)



NASA/Johns Hopkins APL

SPACE

NASA says its asteroid defense test was a success

NASA smashed a spacecraft into an asteroid in an attempt to throw it off course. The mission succeeded beyond expectations, officials said.

"NASA says its mission to knock an asteroid off course — a test of planetary defense — **succeeded beyond its expectations**" (npr.org, emphasis mine)

Data Visualization with `{ggplot2}`

Week 3

Agenda

`{ggplot2}`

- syntax
- continuous data visualizations
- categorical data visualizations
- options
 - color/fill
 - transparency
 - labels
 - facets

Learning Objectives

- Understand the basic syntax requirements for `{ggplot2}`
- Recognize various options for displaying continuous and categorical data
- Familiarity with various `{ggplot2}` options
 - color/fill
 - transparency
 - labels
 - facets

Share!

{datapasta}

- Copy and paste data to and from [R](#)
- VERY handy!
- Good for [reprex](#)
 - posting questions on Rstudio Community or stackoverflow

[demo]





The background of the slide is a dark navy blue field filled with numerous small, multi-colored hexagons and dots. The colors include red, yellow, green, blue, orange, grey, and white, scattered randomly across the entire area. The word "tidyverse" is centered in the middle of the slide in a white, lowercase, sans-serif font.

tidyverse

Providing grammar for:

- Graphics
 - `{ggplot2}`
- Data manipulations
 - `{dplyr}`
 - `{tidyr}`
- Expanding area of specialized topics
 - `{lubridate}`
 - `{glue}`
 - `{tidymodels}`
- Many more...

Providing grammar for:

- Graphics
 - `{ggplot2}`
- Data manipulations
 - `{dplyr}`
 - `{tidyr}`
- Expanding area of specialized topics
 - `{lubridate}`
 - `{glue}`
 - `{tidymodels}`
- Many more...

{ggplot2}

The `{ggplot2}` package

`gg` stands for the "grammar of graphics"

Resources

The `{ggplot2}` package is one of the most popular [R](#) packages, and there are many resources to learn the syntax

- ggplot2 book (email me for digital copy)
- RStudio [cheat sheet](#)
 - Can be helpful, perhaps more so after a little experience
- [R Graphics Cookbook](#)
- [R Graph Gallery](#)
 - past students have really liked this one

Components

Every ggplot plot has three components

1. **data**

- the `data` used to produce the plot

2. **aesthetic mappings (aes)**

- between variables and visual properties

3. **layer(s)**

- usually through the `geom_*()` function to produce geometric shapes to be rendered

`ggplot()` always takes a data frame (tibble) as the first argument

Basic syntax

```
ggplot(data, aes(x = xvar, y = yvar )) +  
  geom_function()
```

`ggplot()` = the function is `ggplot` and the package is `{ggplot2}`

`data` = the data to be plotted

`aes(x = xvar, y = yvar)` = the `aesthetic` mappings

`geom_function()` = the `geometrics` of the plot; the "function" here represents any of the `geom_*` offerings

note the `+` and NOT the `%>%`

{ggplot2} template

```
ggplot(data, aes(mappings)) +  
  geom_function()
```


or equivalently

```
data %>%  
  ggplot(aes(mappings)) +  
  geom_function()
```

Some data for today

penguins from {palmerpenguins}

[run the following]



```
library(tidyverse)
library(palmerpenguins)
```

```
head(penguins)
```


```
## # A tibble: 6 x 8
##   species island    bill_length_mm bill_depth_mm flipper_l~1 body_~2 sex
##   <fct>    <fct>          <dbl>          <dbl>        <int>    <int> <fct>
## 1 Adelie  Torgersen         39.1          18.7         181     3750 male
## 2 Adelie  Torgersen         39.5          17.4         186     3800 fema~
## 3 Adelie  Torgersen         40.3           18         195     3250 fema~
## 4 Adelie  Torgersen          NA           NA           NA         NA <NA>
## 5 Adelie  Torgersen         36.7          19.3         193     3450 fema~
## 6 Adelie  Torgersen         39.3          20.6         190     3650 male
## # ... with abbreviated variable names 1: flipper_length_mm, 2: body_mass_g
```

Continuous Data

Setting up a plot

Run the following code. What do you see?


```
ggplot(penguins, aes(bill_length_mm, body_mass_g))
```



It's ready for you to add some **geometric layers**...what should we add?

How about points?

```
ggplot(penguins, aes(bill_length_mm, body_mass_g)) +  
  geom_point()
```



Adding layers

- In the previous slide, we added a layer of points
- The `geom_point()` layer is a function, complete with it's own arguments

Let's change the color of the points


How would you change the color of the points?

or

How would you find out about changing the color of the points?

color

```
ggplot(penguins, aes(bill_length_mm, body_mass_g)) +  
  geom_point(color = "magenta")
```




]

Add another layer

Let's add a smoothed line with `geom_smooth()`

```
ggplot(penguins, aes(bill_length_mm, body_mass_g)) +  
  geom_point(color = "magenta") +  
  geom_smooth()
```



You try

You probably got the **message** below when you ran (defaults)

```
ggplot(penguins, aes(bill_length_mm, body_mass_g)) +  
  geom_point(color = "magenta") +  
  geom_smooth()
```

`geom_smooth()` using method = 'loess' and formula 'y ~ x'

Change the **method** to "lm"

Let's do this one together

Look at the help page – [?geom_smooth](#)

1. Remove the confidence interval around the line
2. Now change the *SE* band to reflect a 68% confidence interval

color: global vs. conditional

Prior examples changed colors globally

- `geom_point(color = "magenta")`


Use `aes()` to access variables, and color according to a specific variable

- We use variable names within `aes()`

Let's check the data again (`head()`) and the "species" variable (`table()`)

[let's do this together]

```
ggplot(penguins, aes(bill_length_mm, body_mass_g)) +  
  geom_point(aes(color = species))
```



color: global vs. conditional

```
ggplot(penguins, aes(bill_length_mm, body_mass_g)) +  
  geom_point(aes(color = species))
```

- When we did `geom_point(color = "magenta")` we put quotes around the color
- Why now is "species" not in quotes?
 - color names/hex codes are in quotes **NOT** in the `aes()`
 - variable names are in the `aes()` **NOT** in quotes
 - `aes()` is where you map to your data!

Conditional flow through layers

If we use something like `color = "x"` in the first `aesthetic`, it will carry on through all additional layers

- These two codes are the same:

```
ggplot(penguins, aes(bill_length_mm, body_mass_g)) +  
  geom_point(aes(color = species))
```

```
ggplot(penguins, aes(bill_length_mm, body_mass_g, color = species)) +  
  geom_point()
```

- But these two are not...why? [\[run to find out\]](#)

```
ggplot(penguins, aes(bill_length_mm, body_mass_g)) +  
  geom_point(aes(color = species)) +  
  geom_smooth()
```


```
ggplot(penguins, aes(bill_length_mm, body_mass_g, color = species)) +  
  geom_point() +  
  geom_smooth()
```

Be mindful with aes()

Using `aes()` when you **don't** need it

What is happening here?

```
ggplot(penguins, aes(bill_length_mm, body_mass_g)) +  
  geom_point(aes(color = "blue")) +  
  geom_smooth()
```



Be mindful with aes()

Not using `aes()` when you need it

What is happening here?

```
ggplot(penguins, aes(bill_length_mm, body_mass_g)) +  
  geom_point(color = species) +  
  geom_smooth()
```

```
## Error in layer(data = data, mapping = mapping, stat = stat, geom = GeomPoint)
```

Kind of helpful message here.

Themes

Let's talk themes

- The default is `theme_gray`
 - I don't like it
- But there are a lot of build-in alternative in `{ggplot2}`
 - `theme_minimal` is my favorite
- Check out the `{ggthemes}` package for a lot of alternatives
 - These days I nearly always use the `colorblind` theme for discrete values in my plots
- Check out the `{ggthemeassist}` add-in

More themes

- The `{hrbrthemes}` are nice
- Consider [building your own theme!](#)
- Or Google around
- Set the theme globally
 - One of the first lines in your .Rmd file could be:


```
theme_set(theme_minimal())
```

- I often put this in the `setup` chunk


Get a little fancy

- You can use `geom_point()` for more than one layer
- You can also use a different data source on a layer
- Use these two properties to highlight points
 - How about penguins from Torgersen Island?


```
ggplot(penguins, aes(bill_length_mm, body_mass_g)) +  
  geom_point(color = "black")
```




```
ggplot(penguins, aes(bill_length_mm, body_mass_g)) +  
  geom_point(color = "black") +  
  geom_point(data = dplyr::filter(penguins, island == "Torgersen"), color = "magenta")
```



```
ggplot(penguins, aes(bill_length_mm, body_mass_g)) +  
  geom_point(color = "black") +  
  geom_point(data = dplyr::filter(penguins, island == "Torgersen"), color = "magenta") +  
  geom_smooth(se = FALSE)
```




```
ggplot(penguins, aes(bill_length_mm, body_mass_g)) +  
  geom_point(color = "black") +  
  geom_point(data = dplyr::filter(penguins, island == "Torgersen"), color = "magenta") +  
  geom_smooth(se = FALSE) +  
  theme_minimal()
```



Another option

`{gghighlight}` varying flexibility

```
ggplot(penguins, aes(bill_length_mm, body_mass_g)) +  
  geom_point(color = "magenta") +  
  gghighlight::gghighlight(island == "Torgersen") +  
  theme_minimal()
```



Line plots

- When should you use line plots instead of smooths?
 - usually when time is involved
- What are some good candidate data for line plots?
 - observed versus model-implied (estimated)

geom_line()

Classic time series example

economics data from `{ggplot2}`


```
economics
```

```
## # A tibble: 574 x 6
##   date      pce    pop psavert uempmed unemploy
##   <date>    <dbl> <dbl>   <dbl>   <dbl>   <dbl>
## 1 1967-07-01  507. 198712   12.6     4.5    2944
## 2 1967-08-01  510. 198911   12.6     4.7    2945
## 3 1967-09-01  516. 199113   11.9     4.6    2958
## 4 1967-10-01  512. 199311   12.9     4.9    3143
## 5 1967-11-01  517. 199498   12.8     4.7    3066
## 6 1967-12-01  525. 199657   11.8     4.8    3018
## 7 1968-01-01  531. 199808   11.7     5.1    2878
## 8 1968-02-01  534. 199920   12.3     4.5    3001
## 9 1968-03-01  544. 200056   11.7     4.1    2877
## 10 1968-04-01  544 200208   12.3     4.6    2709
## # ... with 564 more rows
```

Let's try it

How do you think we'd fit a line plot to these data, showing unemployment ("unemploy") over time?

```
ggplot(economics, aes(date, unemploy)) +  
  geom_line()
```



Layers

What happens when we layer `geom_line` and `geom_point`?

```
ggplot(economics, aes(date, unemploy)) +  
  geom_line() +  
  geom_point()
```

[try it!]


Not the best instance of this. It would work better on a plot with fewer time points, but you get the idea.

Labels

`"labs()"`


Axis Labels

```
ggplot(economics, aes(date, unemploy)) +  
  geom_line() +  
  theme_minimal() +  
  labs(x = "Date",  
       y = "Unemployment")
```




Title

```
ggplot(economics, aes(date, unemploy)) +  
  geom_line() +  
  theme_minimal() +  
  labs(x = "Date",  
        y = "Unemployment",  
        title = "Unemployment Over Time")
```




Subtitle

```
ggplot(economics, aes(date, unemploy)) +  
  geom_line() +  
  theme_minimal() +  
  labs(x = "Date",  
        y = "Unemployment",  
        title = "Unemployment Over Time",  
        subtitle = "This is the subtitle")
```



Caption

```
ggplot(economics, aes(date, unemploy)) +  
  geom_line() +  
  theme_minimal() +  
  labs(x = "Date",  
       y = "Unemployment",  
       title = "Unemployment Over Time",  
       subtitle = "This is the subtitle",  
       caption = "Created by Joe Nese")
```




Tag

```
ggplot(economics, aes(date, unemploy)) +  
  geom_line() +  
  theme_minimal() +  
  labs(x = "Date",  
        y = "Unemployment",  
        title = "Unemployment Over Time",  
        subtitle = "This is the subtitle",  
        caption = "Created by Joe Nese",  
        tag = "(A)")
```

Legend (one way)

```
ggplot(penguins, aes(bill_length_mm, body_mass_g, color = species)) +  
  geom_point() +  
  theme_minimal() +  
  labs(x = "Bill Length (mm)",  
        y = "Body Mass (g)",  
        color = "Species!") # 'color = '
```




Facets

Faceting

- One of the most powerful features of `{ggplot2}`
- Produce n plots by a specific variable
- `facet_wrap()`
 - wrap a sequence of panels into two dimensions
 - based on variables(s)

Faceting

```
ggplot(penguins, aes(bill_length_mm, body_mass_g)) +  
  geom_point() +  
  geom_smooth() +  
  facet_wrap(~species)
```



Careful about the ~

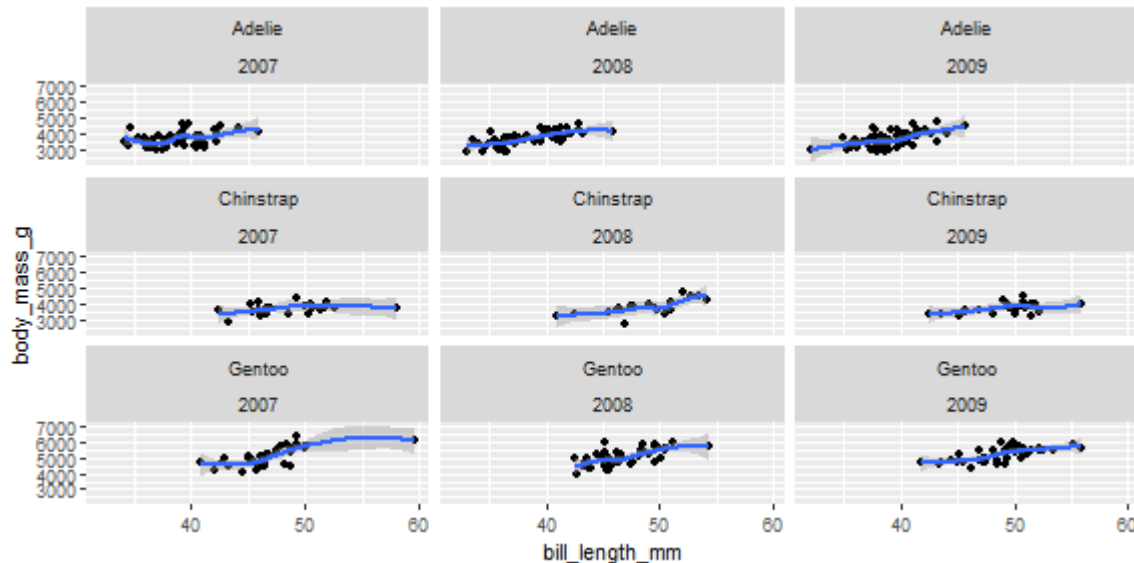
```
ggplot(penguins, aes(bill_length_mm, body_mass_g)) +  
  geom_point() +  
  geom_smooth() +  
  facet_wrap(species)
```

```
### Error in validate_facets(x): object 'species' not found
```


Faceting

two variables (like a matrix)

```
ggplot(penguins, aes(bill_length_mm, body_mass_g)) +  
  geom_point() +  
  geom_smooth() +  
  facet_wrap(species ~ year)
```



Alternative specification - vars()

```
ggplot(penguins, aes(bill_length_mm, body_mass_g)) +  
  geom_point() +  
  geom_smooth() +  
  facet_wrap(vars(species))
```

```
ggplot(penguins, aes(bill_length_mm, body_mass_g)) +  
  geom_point() +  
  geom_smooth() +  
  facet_wrap(vars(species, year))
```

Heatmaps

A heatmap is a literal way of visualizing a table of numbers, where you substitute the numbers with colored cells. -- Nathan Yau

- Useful for finding highs and lows - and sometimes patterns
- They don't always work well


Example with correlations

```
corr <- cor(mtcars)


pc <- corr %>%
  as.data.frame() %>%
  mutate(row = rownames(.)) %>%
  pivot_longer(
    cols = -row,
    names_to = "col",
    values_to = "cor"
  )
head(pc)
```

```
## # A tibble: 6 x 3
##   row    col    cor
##   <chr> <chr>  <dbl>
## 1 mpg   mpg     1
## 2 mpg   cyl   -0.852
## 3 mpg   disp  -0.848
## 4 mpg   hp    -0.776
## 5 mpg   drat   0.681
## 6 mpg   wt    -0.868
```

```
ggplot(pc, aes(row, col, fill = cor)) +  
  geom_tile()
```



```
ggplot(pc, aes(row, col, fill = cor)) +  
  geom_tile() +  
  scale_fill_viridis_c()
```



Categorical Data

Data

`{fivethirtyeight}` package

- `college_grad_students`

```
theme_set(theme_minimal(base_size = 16))

#install.packages("fivethirtyeight")
library(fivethirtyeight)
# View(college_grad_students)
d <- college_grad_students # simpler reference
d
```


```
## # A tibble: 173 x 22
```

	major~1	major	major~2	grad_~3	grad_~4	grad_~5	grad_~6	grad_~7	grad_~8	gr
	<int>	<chr>	<chr>	<int>	<int>	<int>	<int>	<int>	<dbl>	
## 1	5601	Cons~	Indust~	9173	200	7098	6511	681	0.0875	1
## 2	6004	Comm~	Arts	53864	882	40492	29553	2482	0.0578	
## 3	6211	Hosp~	Busine~	24417	437	18368	14784	1465	0.0739	1
## 4	2201	Cosm~	Indust~	5411	72	3590	2701	316	0.0809	
## 5	2001	Comm~	Comput~	9109	171	7512	5622	466	0.0584	
## 6	3201	Cour~	Law & ~	1542	22	1008	860	0	0	1
## 7	6206	Mark~	Busine~	190996	3738	151570	123045	8324	0.0521	1

Histogram

Histogram of "grad_total"

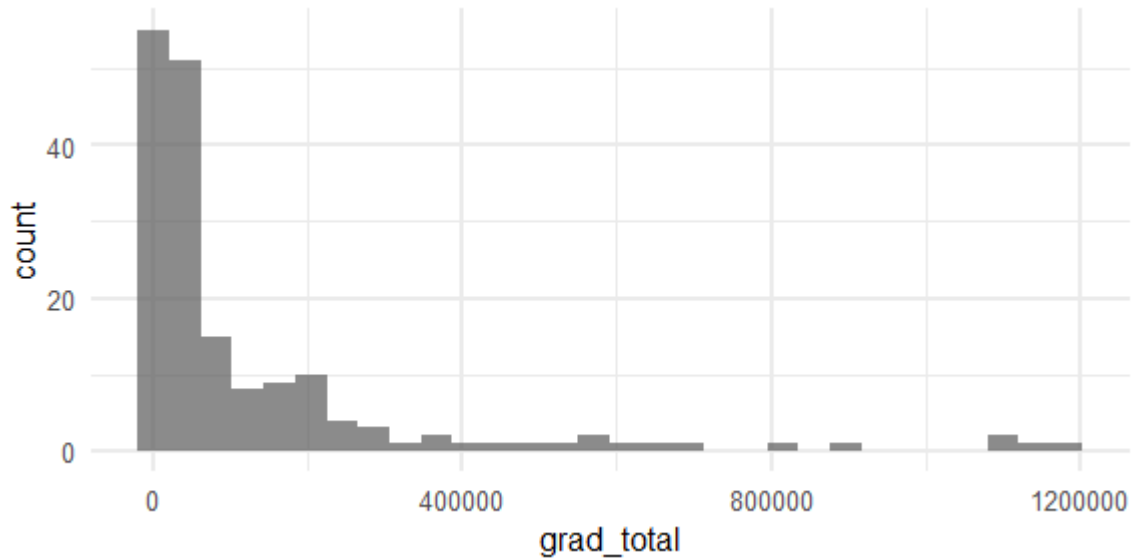
```
ggplot(d, aes(x = grad_total)) +  
  geom_histogram()
```



Transparency - alpha

Add some transparency – perhaps this looks nicer

```
ggplot(d, aes(x = grad_total)) +  
  geom_histogram(alpha = 0.7)
```



color vs. fill

In general

- `color` defines the color a geom is *outlined*
- `fill` defines the color a geom is *filled*

For example:


- `geom_point()` default has only has a color and **NO** fill because they're just points
- Point shapes 21-24 include both a color and a fill

How would we change the color of this plot?

```
ggplot(d, aes(x = grad_total)) +  
  geom_histogram(alpha = 0.7)
```


How would we change the color of this plot?

```
ggplot(d, aes(x = grad_total)) +  
  geom_histogram(alpha = 0.7, color = "magenta")
```



How would we change the color fill of this plot?

```
ggplot(d, aes(x = grad_total)) +  
  geom_histogram(alpha = 0.7, fill = "magenta")
```




Color by variable

What if we wanted different colors by a variable

- *major_category*

```
ggplot(d, aes(x = grad_total)) +  
  geom_histogram(aes(fill = major_category), alpha = 0.7)
```



Density plot


Alternative representation of distribution

- Think of it as a smoothed histogram (uses kernel smoothing)
- The depiction of the distribution is **NOT** determined by the number of the bins you use, as are histograms

Density plot

`geom_density()`


```
ggplot(d, aes(x = grad_total)) +  
  geom_density()
```



Density plot

Now let's fill by *major_category*

```
ggplot(d, aes(x = grad_total)) +  
  geom_density(aes(fill = major_category), alpha = 0.2)
```



Not a good plot, just an example


Possible alternative? `facet_wrap`

```
ggplot(d, aes(x = grad_total)) +  
  geom_density(alpha = 0.2) +  
  facet_wrap(~major_category)
```

Even better

density ridges

```
library(ggribes)  
ggplot(d, aes(x = grad_total, y = major_category)) +  
  geom_density_ridges()
```



Density Ridges

```
ggplot(penguins, aes(bill_length_mm , species)) +  
  geom_density_ridges()
```

fill

```
ggplot(penguins, aes(bill_length_mm , species)) +  
  geom_density_ridges(aes(fill = factor(year)))
```

Add transparency for clarity

```
ggplot(penguins, aes(bill_length_mm , species)) +  
  geom_density_ridges(aes(fill = factor(year)),  
                      alpha = 0.5)
```

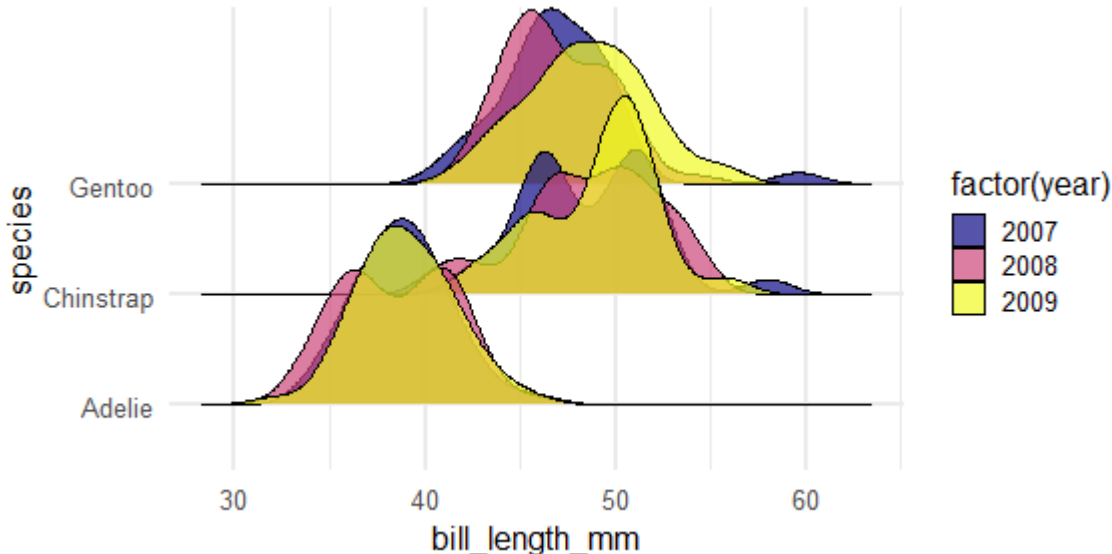
Viridis

- easier to read by those with colorblindness
- prints well in gray scale

```
ggplot(penguins, aes(bill_length_mm , species)) +  
  geom_density_ridges(aes(fill = factor(year)),  
                      alpha = 0.5) +  
  scale_fill_viridis_d()
```


Same fill function, different "option"

```
ggplot(penguins, aes(bill_length_mm , species)) +  
  geom_density_ridges(aes(fill = factor(year)),  
    alpha = 0.7) +  
  scale_fill_viridis_d(option = "plasma")
```



Candy rankings

{fivethirtyeight}

```
candy <- candy_rankings %>%
  pivot_longer(
    cols = chocolate:pluribus,
    names_to = "type",
    values_to = "foo") %>%
  filter(foo) %>%
  select(-foo)
candy
```

```
## # A tibble: 197 x 5
```

	competitorname	sugarpercent	pricepercent	winpercent	type
	<chr>	<dbl>	<dbl>	<dbl>	<chr>
## 1	100 Grand	0.732	0.860	67.0	chocolate
## 2	100 Grand	0.732	0.860	67.0	caramel
## 3	100 Grand	0.732	0.860	67.0	crispedricewafer
## 4	100 Grand	0.732	0.860	67.0	bar
## 5	3 Musketeers	0.604	0.511	67.6	chocolate
## 6	3 Musketeers	0.604	0.511	67.6	nougat
## 7	3 Musketeers	0.604	0.511	67.6	bar
## 8	Air Heads	0.906	0.511	52.3	fruity

Boxplot

```
ggplot(candy, aes(type, sugarpercent)) +  
  geom_boxplot()
```

Violin plots

```
ggplot(candy, aes(type, sugarpercent)) +  
  geom_violin()
```

Bar Charts

```
head(mpg)
```

```
## # A tibble: 6 x 11
##   manufacturer model displ  year   cyl trans      drv    cty   hwy fl    c
##   <chr>          <chr> <dbl> <int> <int> <chr>    <chr> <int> <int> <chr> <
## 1 audi          a4      1.8  1999     4 auto(l5)  f      18    29 p     c
## 2 audi          a4      1.8  1999     4 manual(m5) f      21    29 p     c
## 3 audi          a4      2    2008     4 manual(m6) f      20    31 p     c
## 4 audi          a4      2    2008     4 auto(av)   f      21    30 p     c
## 5 audi          a4      2.8  1999     6 auto(l5)  f      16    26 p     c
## 6 audi          a4      2.8  1999     6 manual(m5) f      18    26 p     c
```

Bar/Col Charts

`geom_bar()`

- expects x **OR** y
- counts rows
- if you want to count the number of cases at each x or y position
- makes the height of the bar proportional to the number of cases in each group
- uses `stat_count()` by default


`geom_col()`

- expects x **AND** y
- expects numbers in your data
- if you want the heights of the bars to represent values in the data
- leaves the data as is
- uses `stat_identity()` by default

geom_bar()

mpg data

```
ggplot(mpg, aes(class)) + # one variable in the 'aes()'
  geom_bar() # counts the rows per class
```



summarized_mpg data

```
summarized_mpg <- mpg %>%  
  group_by(class) %>%  
  count()
```


```
head(summarized_mpg)
```

```
## # A tibble: 6 x 2  
## # Groups:   class [6]  
##   class      n  
##   <chr>    <int>  
## 1 2seater      5  
## 2 compact    47  
## 3 midsize    41  
## 4 minivan    11  
## 5 pickup     33  
## 6 subcompact  35
```


geom_col()


summarized_mpg data

```
ggplot(summarized_mpg, aes(class, n)) + # two variables in the 'aes()'
  geom_col() # data has the rows per class in "n"
```




geom_bar() default

```
ggplot(mpg, aes(class)) +  
  geom_bar()
```




geom_bar(stat = "identity")

```
ggplot(summarized_mpg, aes(class,  
  geom_bar(stat = "identity")
```




geom_bar() default

```
ggplot(mpg, aes(class)) +  
  geom_bar()
```




geom_col() default

```
ggplot(summarized_mpg, aes(class,  
  geom_col())
```



geom_bar() default

```
ggplot(mpg, aes(class)) +  
  geom_bar()
```



geom_bar() uh-oh

```
ggplot(summarized_mpg, aes(class))  
  geom_bar()
```




What happened?

Let's look at our data again

```
summarized_mpg
```

```
## # A tibble: 7 x 2
## # Groups:   class [7]
##   class      n
##   <chr>    <int>
## 1 2seater      5
## 2 compact    47
## 3 midsize    41
## 4 minivan    11
## 5 pickup     33
## 6 subcompact 35
## 7 suv        62
```

```
ggplot(summarized_mpg, aes(class))
  geom_bar()
```




Change the order of the x-axis

- We'll talk about better methods for this when we talk about [factors](#) in Week 9
- Basically - define x-axis variable as a [factor](#) with levels in the order you want

```
(mean_hmiles <- mpg %>%  
  group_by(class) %>%  
  summarize(mean_hwy = mean(hwy)))
```


```
## # A tibble: 7 x 2  
##   class      mean_hwy  
##   <chr>      <dbl>  
## 1 2seater    24.8  
## 2 compact   28.3  
## 3 midsize   27.3  
## 4 minivan   22.4  
## 5 pickup    16.9  
## 6 subcompact 28.1  
## 7 suv       18.1
```

```
mean_hmiles %>%
  ggplot(aes(class, mean_hwy)) +
  geom_col()
```



```
mean_hmiles %>%
  mutate(class = factor(class,
                        levels =
```

```
ggplot(aes(class, mean_hwy)) +
  geom_col()
```



More bar plot options

Data


```
eclsk <- haven::read_sav(here::here("data", "ecls-k_samp.sav")) %>%  
  rio::characterize() %>%  
  janitor::clean_names()  
  
ecls_smry <- eclsk %>%  
  group_by(k_type, ethnic) %>%  
  summarize(t1r_mean = mean(t1rscale))  
  
ecls_smry
```

```
## # A tibble: 16 x 3  
## # Groups:   k_type [2]  
##   k_type    ethnic                                t1r_mean  
##   <chr>    <chr>                                <dbl>  
## 1 full-day AMERICAN INDIAN OR ALASKA NATIVE      17.5  
## 2 full-day ASIAN                                31.8  
## 3 full-day BLACK OR AFRICAN AMERICAN, NON-HISPANIC 21.1  
## 4 full-day HISPANIC, RACE NOT SPECIFIED           20.3  
## 5 full-day HISPANIC, RACE SPECIFIED               22.3  
## 6 full-day MORE THAN ONE RACE, NON HISPANIC       20.8
```


Stacked bar plot


Look for effects in "ethnicity" by "k_type" (full/half day K)

```
ggplot(ecls_smry, aes(ethnic, t1r_mean)) +  
  geom_col(aes(fill = k_type))
```



Grouped bar plot

```
ggplot(ecls_smry, aes(ethnic, t1r_mean)) +  
  geom_col(aes(fill = k_type),  
           position = "dodge")
```



Rotating Labels


I have to look [this](#) up every time

```
ggplot(ecls_smry, aes(ethnic, t1r_mean)) +  
  geom_col(aes(fill = k_type),  
           position = "dodge") +  
  theme(axis.text.x = element_text(angle = 90, hjust = 1))
```

Flip the coordinates

`coord_flip()`


```
ggplot(ecls_smry, aes(ethnic, t1r_mean)) +  
  geom_col(aes(fill = k_type),  
           position = "dodge") +  
  theme(axis.text.x = element_text(angle = 90, hjust = 1)) +  
  coord_flip()
```



Alternatively


`facet_wrap()`

```
ggplot(ecls_smry, aes(k_type, t1r_mean)) +  
  geom_col(alpha = 0.8) +  
  facet_wrap(~ethnic)
```



Sometimes some redundancy works well

```
ggplot(ecls_smry, aes(k_type, t1r_mean, fill = k_type)) +  
  geom_col(alpha = 0.8) +  
  facet_wrap(~ethnic)
```



geom_*() Review

- `geom_point()`
- `geom_smooth()`
- `geom_line()`
- `geom_tile()`
- `geom_histogram()`
- `geom_density()`
- `ggribes::geom_density_ridges()`
- `geom_boxplot()`
- `geom_violin()`
- `geom_bar()`
- `geom_col()`

Challenge

- Start a new R project
- Create a new script, save it as "practice-plots.R"
- Load the `{tidyverse}`
- Print the `msleep` dataset to see its structure (it's from `{ggplot2}`)

For each of the following, produce a separate plot

1. Plot the relation between "sleep_total" and "brainwt" (with "brainwt" as the DV) - scatter plot
2. Overlay a smooth on the previous plot
3. Color the points by "vore", but fit a single smooth
4. Fit separate smooths by "vore", but with all points being gray
5. Omit the standard error of the smooths
6. Use `ylim()` as an additional layer to restrict the y-axis to range from 0 to 5

Next time

Before next class

- Reading
 - [R4DS 5](#)
- Supplemental Learning
 - [RStudio Primer: Working with Tibbles](#)
 - [Rbootcamp: Ch 3](#)
 - [Codecademy: Aggregates in R](#)
- Homework
 - **Homework 2**
 - **Homework 3**


```
renderthis::to_pdf(here::here("slides", "w3_ggplot2.Rmd"), complex_slide
```